

SENTIMENTAL ANALYSIS OF IMDB MOVIE REVIEWS

A MINI PROJECT REPORT

Submitted by

AISWARYA SREEKUMAR(MZW21CS002)

HARIHARAN R(MZW21CS008)

SOUMYA S NAIR(MZW21CS016)

to

the APJ Abdul Kalam Technological University

in partial fulfilment of the requirement for the award of the Degree

of

Bachelor of Technology in Computer Science and Engineering



MZIST

Department of Computer Science and Engineering

MOUNT ZION INSTITUTE OF SCIENCE AND TECHNOLOGY

KOZHUVALLOOR, CHENGANNUR

MAY 2024

DECLARATION

We undersigned hereby declare that the mini project report "Sentimental Analysis of IMDb Movie Reviews", submitted for partial fulfilment of the requirements for the award of degree of Bachelor of Technology of the APJ Abdul Kalam Technological University, Kerala is a bonafide work done by us under supervision of Assistant Professor **Mrs. Hazeena Khalid**, Department of Computer Science and Engineering, Mount Zion Institute of Science and Technology. This submission represents our ideas in our own words and where ideas or words of others have been included, we have adequately and accurately cited and referenced the original sources. We also declare that we have adhered to ethics of academic honesty and integrity and have not misrepresented or fabricated any data or idea or fact or source in our submission. We understand that any violation of the above will be a cause for disciplinary action by the Institute and University and can also evoke penal action from the sources which have thus not been properly cited or from whom proper permission has not been obtained. This report has not been previously formed the basis for the award of any degree, diploma or similar title of any other University.

Kozhuvalloor

Aiswarya Sreekumar

Hariharan R

Soumya S Nair

**DEPARTMENT OF COMPUTER SCIENCE AND
ENGINEERING
MOUNT ZION INSTITUTE OF SCIENCE AND
TECHNOLOGY KOZHUVALLOOR**



MZIST

CERTIFICATE

This is to certify that the report entitled **SENTIMENTAL ANALYSIS OF IMDb MOVIE REVIEWS** submitted by **AISWARYA SREEKUMAR (MZW21CS002), HARIHARAN R (MZW21CS008), SOUMYA S NAIR (MZW21CS016)** to the APJ Abdul Kalam Technological University in partial fulfilment of the requirements for the award of the Degree of Bachelor of Technology in Computer Science and Engineering is a bonafide record of the mini project work CSD334 carried out by them under our guidance and supervision at MOUNT ZION INSTITUTE OF SCIENCE AND TECHNOLOGY, KOZHUVALLOOR during the year of 2021-2025. This report in any form has not been submitted to any other University or Institute for any purpose.

Internal Supervisor

External Supervisor

Internal Guide
Mrs. Hazeena Khalid
Assistant Professor
Dept. of Computer
Science & Engg.

Project Co-ordinator
Mrs. Thara TR
Assistant Professor
Dept. of Computer
Science & Engg.

Head of the Department
Mr. Jacob Joseph
Assistant Professor
Dept. of Computer
Science & Engg.

ACKNOWLEDGEMENT

We are grateful to the co-operation and constant encouragement by our Head of the Department **Mr. JACOB JOSEPH** and coordinator **Mrs. THARA T R.** Their regular suggestions made our work easy and proficient. We should like to express our profound gratitude to our guide **Mrs. HAZEENA KHALID** for her invaluable support, encouragement, supervision and useful suggestions throughout this project work. Their moral support and continuous guidance enabled us to complete our work successfully. We wish to express our appreciation to all the staff members of CSE department of our college who helped to overcome our doubts in doing this project.

We are heavily indebted to our Principal **Prof. Dr. K MATHEW** for his constant inspiration assistance throughout the project. We wish to thank our parents for their undivided support and interest who inspired us and encouraged us to go our own way.

ABSTRACT

Sentiment Analysis has been a classic field of research in Natural Language Processing, Text Analysis and Linguistics. It essentially attempts to identify, categorize and possibly quantify, the opinions expressed in a piece of text and determine the author's attitude toward a topic, product or situation. This has widespread application in Recommender systems for predicting the preferences of users and in e-commerce websites to analyze customer feedback and reviews. Based on the sentiments extracted from the data, companies can better understand their customers and align their businesses accordingly. The sentiment analysis is an emerging research area where vast amount of data is being analyzed, to generate useful insights in regards to a specific topic. It is an effective tool which can serve governments, corporations and even consumers. Text emotion recognizing lays a key role in this framework. Researchers in the fields of natural language processing (NLP) and machine learning (ML) have explored a variety of methods to implement the process with highest accuracy possible. In this paper the Long Short-Term Memory (LSTM) and Convolutional Neural Network (CNN) classifier are used for analyzing sentiments of the IMDb movie reviews. It is based on the Recurrent Neural Network (RNN) algorithm. The data is effectively preprocessed and partitioned to enhance the post classification performance. The classification performance is studied in terms of accuracy. Results show a best classification accuracy of 89.9.

TABLE OF CONTENTS

CONTENTS	PAGE NO.
ACKNOWLEDGEMENT	i
ABSTRACT	ii
LIST OF FIGURES	v
ABBREVIATIONS	vi
1.INTRODUCTION	
1.1 GENERAL BACKGROUND	1
1.2 OBJECTIVE	2
1.3 SCOPE	2
1.4 SCHEME	3
2.LITERATURE REVIEW	4
3.METHODOLOGY	
3.1 DATA COLLECTION	14
3.2 DATA PREPROCESSING	17
3.3 PROJECT PHASE	18
4.SYSTEM DESIGN	
4.1ARCHITECTURE	
4.1.1 Architecture of LSTM	23
4.1.2 Architecture of CNN	29
4.2 DESIGN	
4.2.1 Data Collection and Processing	33
4.2.2 Model Training and Evaluation	34

4.2.2.1 Model Building	35
4.2.2.2 Model Training	36
4.2.2.3 Model Evaluation	36
4.2.3 Integration and Deployment	37
5.RESULTS	39
6.CONCLUSIONS	
6.1 CONCLUSIONS	48
6.2 FUTURE WORK	48
7.REFERENCES	50

LIST OF FIGURES

Figure No	Title	Page No
Fig 1	Architecture of LSTM	23
Fig 2	Architecture of CNN	29
Fig 3	Data Collection and Preprocessing	33
Fig 4	Model Training and Evaluation	35
Fig 5	Integration and Deployment	37
Fig 6	Testing Dataset	39
Fig 7	Training Dataset	40
Fig 8	Training data of CNN Model	41
Fig 9	Testing Data of CNN Model	41
Fig 10	Training Data of LSTM Model	42
Fig 11	Testing Data of LSTM Model	42
Fig 12	CNN Model	43
Fig 13	CNN Model After Building	43
Fig 14	LSTM Model	44
Fig 15	LSTM Model After Building	44
Fig 16	CNN Model	45
Fig 17	CNN Model After Training	45
Fig 18	LSTM Model	46
Fig 19	LSTM Model After Training	46
Fig 20	Web Interface	47
Fig 21	Prediction of Movie Review	47

ABBREVIATIONS

No	Acronym	Meaning
1	LSTM	Long Short Term Memory
2	CNN	Convolutional Neural Network
3	ML	Machine Learning
4	SA	Sentiment Analysis
5	IMDb	Internet Movie Database
6	DL	Deep Learning
7	NLP	Natural Language Processing
8	Adam	Adaptive Moment Estimation

CHAPTER 1

INTRODUCTION

1.1 GENERAL BACKGROUND:

In the age of digital media, platforms like IMDb have become invaluable sources for understanding public sentiment. Sentiment analysis, a branch of NLP, systematically categorizes text based on expressed emotions. By employing machine learning and deep learning, sentiment analysis algorithms can classify text as positive, negative, or neutral, providing insights into sentiment trends. Applied across domains like customer feedback and product reviews, sentiment analysis offers valuable insights into public opinion and brand perception.

In this project, focus on sentiment analysis applied to IMDb movie reviews. The aim of the project is to leverage advanced deep learning architectures, such as LSTM and CNN, for accurate sentiment analysis. LSTM excels in capturing temporal dependencies, making it effective for analyzing sequential data like movie reviews. Meanwhile, CNN, originally for image processing, has been adapted successfully for text analysis. By integrating LSTM and CNN, it seeks to enhance sentiment classification and provide actionable insights into user sentiments towards movies.

The approach involves integrating LSTM and CNN architectures into a unified framework for IMDb movie review sentiment analysis. By combining the strengths of both models, aim to enhance accuracy and robustness in sentiment classification. Additionally, optimization techniques like the Adam optimizer ensure optimal performance in sentiment analysis tasks. Through this project, it contributes to advancing sentiment analysis methodologies, empowering filmmakers and producers with valuable insights for decision-making and marketing strategies.

The Large Movie Review Dataset integrated with Keras is a comprehensive resource comprising intensely polar movie reviews for training and an additional reviews for testing purposes. The dataset is meticulously curated to prevent the presence of more than 30 reviews for any individual movie, ensuring a balanced representation across films. Furthermore, it meticulously maintains an equal distribution of positive and negative reviews within both the training and test sets, contributing to the robustness and fairness of the dataset for sentiment analysis tasks.

In addition to the integration of LSTM and CNN architectures, the project also emphasizes

the importance of pre-processing techniques such as tokenization and word embedding. These techniques enable the models to effectively understand and represent the semantic meaning of words within the context of movie reviews. Furthermore, the project explores the utilization of attention mechanisms to focus on the most relevant parts of the reviews, enhancing the models' interpretability and performance. Additionally, the incorporation of ensemble learning methods could further boost the robustness and generalization capabilities of the sentiment analysis system. Lastly, the project aims to investigate the impact of domain adaptation techniques to tailor the models specifically for IMDB movie reviews, potentially improving their accuracy and applicability in real-world scenarios.

1.2 OBJECTIVE:

The primary objective of this project is to harness the power of advanced deep learning architectures, particularly Long Short-Term Memory (LSTM) and Convolutional Neural Network to conduct accurate sentiment analysis on IMDB movie reviews. By leveraging these sophisticated models, we aim to delve deeper into the nuanced sentiments expressed by users towards various films listed on IMDB. Additionally, our goal is to explore how the integration of LSTM and CNN architectures can enhance the granularity and accuracy of sentiment classification, thereby providing filmmakers, producers, and industry stakeholders with more comprehensive insights into audience perceptions and preferences. Through the implementation of state-of-the-art deep learning techniques, we strive to not only decipher the polarity of reviews but also uncover underlying patterns and trends in audience sentiment, facilitating more informed decision-making processes within the dynamic landscape of the entertainment industry. Certainly! Here's an elaboration of the scope with additional sentences:

1.3 SCOPE:

This project specifically targets sentiment analysis within the context of IMDB movie reviews, aiming to extract nuanced insights into user sentiments towards different films listed on the platform. By focusing on IMDB, which serves as a prominent hub for film enthusiasts and critics alike, we aim to capture a diverse range of opinions and sentiments expressed by users from various demographics and backgrounds. The scope extends beyond merely classifying reviews as positive, negative, or neutral; instead, we aim to

delve deeper into the subtleties of user emotions and perceptions, identifying specific aspects of films that evoke strong reactions or resonate positively with audiences. Furthermore, the project aims to explore the temporal dynamics of sentiment within the context of movie reviews, analyzing how user perceptions may evolve over time in response to factors such as release dates, marketing campaigns, and critical reception. By integrating L S T M and CNN architectures, we seek to not only enhance the accuracy of sentiment classification but also uncover deeper insights into the underlying factors driving audience engagement and satisfaction with films. Ultimately, the scope of this project encompasses a holistic understanding of user sentiments towards IMDb movie reviews, with the overarching goal of providing actionable insights for filmmakers, producers, and industry stakeholders to inform decision-making processes and marketing strategies.

1.4 SCHEME:

The proposed approach involves integrating L S T M and CNN architectures into a unified frame- work tailored for sentiment analysis of IMDb movie reviews. Through this integration, we aim to capitalize on the strengths of both models, leveraging LSTM's ability to capture temporal dependencies and CNN's effectiveness in text analysis. Additionally, we plan to employ optimization techniques such as the Adam optimizer to ensure optimal performance in sentiment analysis tasks. Furthermore, the project emphasizes the significance of pre-processing techniques like tokenization and word embedding, which enable the models to interpret and represent the semantic meaning of words within the context of movie reviews. The utilization of attention mechanisms will allow the models to focus on the most relevant aspects of the reviews, thereby enhancing interpretability and performance. Furthermore, we will explore ensemble learning methods to enhance the robustness and generalization capabilities of the sentiment analysis system. Finally, we aim to investigate do- main adaptation techniques to fine-tune the models specifically for IMDb movie reviews, potentially improving their accuracy and applicability in real-world scenarios.

CHAPTER 2

LITERATURE REVIEW

[1] The work presented in 'Sentiment Analysis on Twitter Data Using Deep Learning approach (2020)' focuses on leveraging data from social networking sites, specifically targeting global product reviews and customer opinions. The proposed system aims to analyze sentiments and classify reviews as positive or negative using a CNN- LSTM based deep learning method. By employing this approach, the system offers automatic feature extraction, resulting in performance improvement and the ability to handle unstructured data effectively. Furthermore, the method facilitates semantic information capture, enhancing the accuracy and depth of sentiment analysis on Twitter data.

Social media platforms, including Twitter, Facebook, and Instagram, have become hubs for users to express their opinions, share experiences, and review products. This user-generated content provides a wealth of information that can offer valuable insights into consumer sentiments, preferences, and behaviors. Sentiment analysis, a branch of natural language processing (NLP), involves analyzing this data to determine the emotional tone or attitude expressed within text, such as whether a product review is positive, negative, or neutral. By leveraging sentiment analysis techniques, businesses can gain a deeper understanding of customer sentiment towards their products or services, enabling them to make informed decisions regarding marketing strategies, product development, and customer engagement.

Despite the wealth of information available on social media platforms, analyzing this data presents several challenges. Unlike structured data found in databases or spreadsheets, social media data is often unstructured, meaning it lacks a predefined data model or organization. Social media posts frequently contain informal language, abbreviations, slang, misspellings, and emojis, making them difficult to interpret using traditional analysis methods. As a result, sentiment analysis on social media data requires advanced techniques capable of understanding and extracting meaning from this unstructured textual data.

Deep learning, a subset of machine learning that involves artificial neural networks with multiple layers, has emerged as a powerful approach for sentiment analysis on social media data. Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs), such as Long Short-Term Memory networks (LSTMs), are commonly used architectures in deep learning for natural language processing tasks. These neural networks excel at learning intricate patterns and relationships within text data, allowing them to effectively capture the nuanced semantics and context present in social media posts. As a result, deep learning models have demonstrated superior performance in sentiment analysis tasks

compared to traditional machine learning approaches.

Before applying deep learning techniques to social media data, preprocessing is typically performed to clean and standardize the text. This preprocessing may involve steps such as tokenization (breaking text into individual words or tokens), stemming (reducing words to their root form), and removing stop words (commonly occurring words with little semantic value). Additionally, word embeddings are often used to represent words in a continuous vector space, capturing semantic relationships and contextual information. Word embedding models, such as Word2Vec and GloVe, learn distributed representations of words based on their surrounding context in large text corpora, enabling deep learning models to effectively capture semantic meaning during sentiment analysis tasks.

The project introduces a novel approach for sentiment analysis on social media data, combining Convolutional Neural Networks (CNNs) and Long Short-Term Memory networks (LSTMs) within a deep learning framework. CNNs are well-suited for capturing local features and patterns within text, while LSTMs excel at modeling sequential dependencies and long-range dependencies in data. By integrating these two architectures, the proposed approach aims to leverage the strengths of both CNNs and LSTMs for sentiment analysis tasks. Additionally, the use of pre-trained word embeddings further enhances the model's ability to automatically extract meaningful features from social media text data, contributing to improved sentiment classification performance.

Finally, this project discusses the evaluation of the proposed CNN- LSTM based approach for sentiment analysis on social media data. Performance evaluation involves testing the model on benchmark datasets and assessing its accuracy, precision, recall, and F1-score, among other metrics. Additionally, the performance of the CNN- LSTM based approach is compared against baseline machine learning methods to validate its effectiveness. Comparative analysis provides insights into the superiority of deep learning approaches for sentiment analysis tasks on social media data, highlighting the potential of the proposed approach to outperform traditional methods.

[2] The work presented in 'Sentiment Analysis using deep learning for use in recommendation systems of various public media applications (2022)' aims to enhance the accuracy of recommendation systems by integrating sentiment analysis to gain deeper insights into user emotions and opinions. The proposed system leverages deep learning algorithms along with recommendation system algorithms to achieve this goal. By adopting this approach, the system offers several advantages including enhanced user experience through personalized recommendations, automatic feature extraction for improved performance, and customization options tailored to individual user preferences.

Traditional methods of sentiment analysis involve several approaches. Lexicon-based methods utilize sentiment dictionaries containing words or phrases with associated sentiment scores. These scores are used to determine the overall sentiment of a piece of text based on the presence and polarity of sentiment-bearing words. Machine learning algorithms, on the other hand, learn patterns and relationships within labeled datasets to classify text into sentiment categories (positive, negative, neutral). Examples include Support Vector Machines (SVM), Naive Bayes classifiers, and decision trees. Rule-based methods employ predefined rules or patterns to identify sentiment cues and expressions in text. While these approaches are effective to some extent, they may struggle with complex language structures, sarcasm, and context-dependent sentiments.

Deep learning has emerged as a powerful paradigm for sentiment analysis, capable of automatically learning intricate patterns and representations from raw text data. Convolutional Neural Networks (CNNs) are effective at capturing local features and patterns within text, making them suitable for tasks like sentiment classification. Recurrent Neural Networks (RNNs), including variants like Long Short-Term Memory networks (LSTMs) and Gated Recurrent Units (GRUs), are well-suited for modeling sequential dependencies in text data. These deep learning architectures can effectively capture the context and semantics of language, leading to improved sentiment analysis performance.

Sentiment analysis plays a vital role in recommendation systems by providing insights into user preferences and opinions. By analyzing sentiment-rich data from various public media sources such as product reviews, forum discussions, blogs, and social media posts, recommendation systems can personalize content and recommendations for users. Deep learning-based sentiment analysis enables these systems to process large volumes of text data efficiently and accurately, resulting in more personalized and relevant recommendations tailored to individual user preferences and sentiments.

Deep learning-based sentiment analysis poses several challenges, including the need for large annotated datasets for model training, ensuring the interpretability and explainability of results, and addressing ethical considerations such as privacy and bias. However, advancements in deep learning techniques, coupled with the abundance of text data available from public media sources, present numerous opportunities for enhancing recommendation system accuracy and effectiveness. Future research in this area may focus on developing novel deep learning architectures, handling multi-lingual and multimodal data, and addressing ethical concerns to further improve recommendation system performance.

Future research directions in sentiment analysis for recommendation systems may involve exploring novel deep learning architectures tailored to specific application domains, such as e-commerce, social media, or entertainment. Additionally, techniques for handling multilingual and multimodal data could

further enhance recommendation system capabilities. Ethical considerations, including privacy preservation and mitigation of algorithmic bias, will continue to be important areas of focus. Integrating sentiment analysis with other recommendation system components, such as collaborative filtering and content-based filtering, could lead to more comprehensive and effective recommendation strategies in various public media applications.

[3] The work presented in 'Dynamic Word Vector Based Sentiment Analysis for Microblog Short Text (2022)' introduces a system designed to recognize emotions in microblog comments, enabling effective analysis and monitoring of social public opinions. The proposed system utilizes a dynamic word vector model combined with a Bidirectional G RU Neural Network to achieve this goal. By employing these methods, the system offers several advantages, including improved accuracy in sentiment analysis, dynamic word vectors for capturing nuanced meanings, the effectiveness of a Bidirectional G RU network in processing sequential data, and the applicability of the approach to social media platforms.

In recent years, sentiment analysis on microblog short texts has become increasingly important due to the widespread use of social media platforms like Twitter and Weibo. Traditional sentiment analysis methods often rely on static word vector models, which may struggle with the polysemy of words—words having multiple meanings depending on context. To address this limitation, researchers have been exploring dynamic word vector models combined with neural network architectures to improve sentiment analysis accuracy and effectiveness.

Static word vector models, such as Word2Vec and GloVe, represent words as fixed vectors in a high-dimensional space based on co-occurrence statistics in large text corpora. However, these models fail to capture the dynamic nature of word meanings, particularly in contexts where words exhibit multiple senses or polysemy. As a result, sentiment analysis performance may be suboptimal when dealing with microblog short texts containing ambiguous or context-dependent language.

Dynamic word vector models aim to address the limitations of static models by capturing contextual variations in word meanings. Models such as ELMo (Embeddings from Language Models) and E R N I E (Enhanced Representation through knowledge Integration) leverage contextual information from pre-trained language models to generate word embeddings that adapt to the surrounding context. By incorporating contextual information, dynamic word vector models can better capture the nuances of language and improve sentiment analysis performance on microblog short texts.

Neural network architectures, particularly recurrent neural networks (RNNs) and variants like

Gated Recurrent Units (GRUs) and Long Short-Term Memory networks (LSTMs), have been widely used for sentiment analysis tasks. These architectures are well-suited for processing sequential data, making them suitable for analyzing microblog short texts. Bidirectional RNNs, in particular, have the advantage of capturing both past and future contexts of words, enhancing the model's ability to understand the semantics of microblog text.

Attention mechanisms have proven effective in improving the performance of neural network models by focusing on relevant parts of the input sequence. By assigning weights to different parts of the input sequence, attention mechanisms enable the model to selectively attend to important information while filtering out noise. Incorporating attention mechanisms into neural network architectures for sentiment analysis can enhance the model's ability to identify keywords and phrases relevant to sentiment expression in microblog short texts.

The proposed model in the project combines dynamic word vector representations obtained from the ERNIE pre-training model with a bidirectional GRU neural network and attention mechanism. The dynamic word vectors capture contextual variations in word meanings, while the bidirectional GRU neural network processes the sequential nature of microblog short texts. The attention mechanism weights keywords in the input sequence, enabling the model to focus on salient information relevant to sentiment analysis. Finally, sentiment classification is performed using a SoftMax layer. The incorporation of a bidirectional GRU neural network enhances the model's ability to process the sequential nature of microblog text data. Bidirectional GRUs capture both past and future contexts of words, enabling the model to extract relevant features and patterns from the input sequence more effectively.

Furthermore, the integration of an attention mechanism enables the model to focus on salient information within the input sequence. By assigning weights to different parts of the text, the attention mechanism allows the model to prioritize keywords and phrases that are most relevant to sentiment expression. This selective attention mechanism helps filter out noise and irrelevant information, leading to improved sentiment analysis performance.

Overall, the high accuracy achieved by the proposed model underscores the importance of dynamic word vectors, neural network architectures, and attention mechanisms in addressing the challenges of sentiment analysis on social media platforms like Weibo. The experimental results provide strong support for the effectiveness of the proposed approach in accurately identifying and classifying sentiment in microblog short texts, thereby contributing to the advancement of sentiment analysis techniques in the context of social media data.

[4] The work presented in 'Sentiment-Based Deep Auto Encoder Recommender System (DAERS(2020))'

introduces a novel approach to address Cold Start Recommendations (CSR) for new users by proposing a low-rank Deep Autoencoder model. This model leverages Deep Autoencoder (DAE) techniques and cross-domain recommendations to effectively handle the challenges of recommending items to users with limited historical data. By incorporating user sentiments, the system offers advantages such as improved generalization, utilization of user sentiments for personalized recommendations, and the efficiency of a low-rank Deep Auto encoder architecture.

Social media platforms like Twitter have become invaluable sources for understanding user sentiment. Sentiment analysis techniques aim to extract and classify sentiments expressed in user-generated content, such as tweets, comments, and reviews. These techniques are crucial for various applications, including market analysis, opinion mining, and recommendation systems. Researchers have developed a wide range of methods, from traditional lexicon-based approaches to more sophisticated deep learning models, to analyze sentiment in social media data effectively.

Cross-domain recommendations involve leveraging user-generated content from different domains to make personalized recommendations. This process requires analyzing user reviews and comments to uncover latent thematic structures in both the source and target domains, even when there is no overlap between users and items across domains. By understanding user sentiments across different domains, recommendation systems can provide tailored recommendations to users based on their preferences and emotional responses, enhancing user satisfaction and engagement.

Data sparsity and the cold start problem are significant challenges for recommendation systems, particularly for new users with limited interaction history. Deep learning techniques offer promising solutions to these challenges by leveraging user sentiments extracted from textual data. The proposed approach integrates deep learning methods, specifically a low-rank Deep Autoencoder (DAE) , to infer user ratings based on sentiment analysis of user reviews. By mapping user behavior to attribute space and reconstructing behavior space from attribute space, the D A E addresses data sparsity and C S R issues effectively, improving recommendation accuracy for both new and existing users.

The proposed approach addresses three key challenges: domain shift, cold start recommendations, and spurious correlations. Domain shift refers to the differences between source and target domains, which can affect the generalization ability of recommendation models. Cold start recommendations involve providing accurate recommendations for new users with limited historical data. Spurious correlations may

arise when unrelated factors influence recommendations. The proposed approach tackles these challenges through the use of deep auto encoders and normalization techniques like Dropout to prevent overfitting, ensuring robust performance across diverse datasets and domains.

Deep auto encoders offer several advantages over traditional linear models in recommendation systems. These models can capture complex nonlinear relationships in the data, allowing them to represent high dimensional feature spaces more effectively. Additionally, techniques like Dropout regularization help prevent overfitting by randomly dropping units during training, improving the model's ability to generalize to unseen data. By leveraging deep auto encoders and sentiment analysis, the proposed approach enhances recommendation accuracy and addresses key challenges in recommendation systems, ultimately improving user satisfaction and engagement.

[5] The work presented in "Sentiment Analysis on User Feedback of a Social Media Platform" suggests a study focused on analyzing the sentiments expressed in user feedback on a social media platform. This entails examining the opinions, emotions, and attitudes conveyed by users through their comments, reviews, or interactions on the platform. The goal is to gain insights into how users perceive and feel about the platform, its features, content, or services. Through sentiment analysis techniques, researchers aim to categorize user feedback as positive, negative, or neutral, providing valuable insights for platform improvement, user engagement enhancement, and decision-making processes.

Sentiment analysis encompasses a variety of methodologies, ranging from traditional machine learning to advanced deep learning models. Basic machine learning methods typically involve pre-processing text data, extracting features, and training classifiers using labeled datasets. Techniques such as bag-of-words, n-grams, and supervised learning algorithms like Support Vector Machines (SVM) or Naive Bayes classifiers are commonly used. In contrast, deep learning models, such as BERT, utilize neural networks to learn complex patterns and relationships in text data. BERT, in particular, employs bidirectional transformers to capture contextual information and semantic meaning from input text, leading to more accurate sentiment analysis results.

Sentiment analysis plays a crucial role in understanding user feedback on social media platforms, which serve as hubs for user-generated content. Organizations can analyze user feedback, comments, reviews, and interactions to gain insights into user sentiments towards their products, services, or platform features. By leveraging sentiment analysis techniques, organizations can identify strengths and weaknesses, address user concerns, and tailor their strategies to enhance user satisfaction and engagement. Additionally, sentiment analysis provides valuable insights into user preferences, trends, and behavior, guiding strategic decision-making processes. Comparative studies have shown that deep learning

models, particularly BERT , often outperform traditional machine learning approaches in sentiment analysis tasks. Deep learning models have the advantage of capturing complex relationships and nuances in language, allowing them to achieve higher accuracy and performance in sentiment classification. While traditional machine learning methods are effective in certain scenarios, deep learning models excel in capturing contextual information and handling large-scale datasets. However, the choice of methodology depends on factors such as the availability of labeled data, computational resources, and the specific requirements of the application.

The results obtained from sentiment analysis have a significant impact on organizational decision-making processes. By understanding user sentiments expressed in feedback, organizations can make informed decisions to improve their products, services, or platform features. Sentiment analysis enables organizations to prioritize areas for improvement, allocate resources effectively, and tailor their strategies to meet user needs and preferences. Additionally, sentiment analysis helps organizations build brand loyalty, maintain a positive reputation, and stay competitive in the dynamic landscape of social media platforms.

Future research in sentiment analysis is focused on addressing challenges and advancing methodologies to enhance accuracy and generalization capabilities. Challenges include improving the interpretability and explainability of deep learning models, mitigating biases in sentiment analysis algorithms, and handling multilingual and multimodal data. Advancements in natural language processing techniques, such as transfer learning and domain adaptation, hold promise for improving the robustness and scalability of sentiment analysis models in diverse social media settings. Continued research and development in these areas will drive innovation and improve the effectiveness of sentiment analysis in informing organizational decision-making processes.

[6] The work presented in "Sentiment Analysis in Social Media: Handling Noisy Data and Detecting Sarcasm Using a Deep Learning Approach" suggests a study focused on analyzing sentiments expressed in social media content. Specifically, it addresses challenges related to noisy data and the detection of sarcasm, leveraging deep learning techniques for improved accuracy. This research likely involves developing methods to preprocess noisy social media data and implement deep learning models capable of detecting sarcasm, ultimately enhancing sentiment analysis in social media contexts.

Deep learning techniques, particularly Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks, have gained popularity in sentiment analysis tasks due to their ability to automatically learn representations from data. These neural network architectures are well-suited for capturing sequential dependencies and contextual information in text data, making them effective

for sentiment analysis in social media content. Researchers have explored various deep learning models and methodologies to enhance sentiment analysis accuracy, especially in the context of noisy and informal social media data.

Social media data often contain noisy elements such as slang, abbreviations, spelling errors, and informal language, which can pose challenges for sentiment analysis. To address this, researchers have developed preprocessing techniques, lexical resources, and manual annotations to clean and preprocess noisy data effectively. Deep learning models, with their capacity to learn from large-scale data, have shown promise in mitigating the impact of noise on sentiment analysis outcomes by automatically learning relevant features and patterns from noisy data.

Detecting sarcasm in social media poses a significant challenge for sentiment analysis due to its nuanced nature. Sarcasm involves expressing sentiments that are opposite or different from the literal meaning of the text. To accurately identify sarcasm, models need to understand the contextual cues, tone, and linguistic patterns present in the text.

Deep learning approaches, particularly Long Short-Term Memory (LSTM) networks, have been increasingly employed for sarcasm detection tasks in social media data. LSTM networks are well-suited for capturing sequential dependencies and contextual information in text data, making them effective for detecting sarcasm, which often relies on subtle linguistic cues and context.

Researchers have explored various features and mechanisms to improve sarcasm detection accuracy. Word embeddings, which represent words in a continuous vector space, capture semantic relationships between words and can help models understand the context in which sarcasm occurs. Context-aware representations consider the broader context of the conversation or the user's history to better interpret the intended meaning of the text. Attention mechanisms focus on relevant parts of the input text, allowing the model to prioritize important information for sarcasm detection.

By leveraging these features and mechanisms, researchers aim to enhance the ability of deep learning models to accurately identify sarcasm in social media content. Continued advancements in deep learning techniques and model architectures hold promise for further improving sarcasm detection accuracy and enabling more robust sentiment analysis in social media contexts.

Performance evaluation metrics such as accuracy, precision, recall, and F1 score are commonly used to assess the effectiveness of sentiment analysis and sarcasm detection models. Accuracy measures the overall correctness of predictions, while precision and recall evaluate the model's ability to correctly classify positive and negative instances. F1 score provides a balance between precision and recall.

Comparative studies and benchmarking against existing methods help validate the performance of deep learning models and identify areas for improvement, guiding the development of more accurate and robust sentiment analysis and sarcasm detection systems.

While deep learning has shown promising results in sentiment analysis and sarcasm detection, several challenges remain. Generalizing models across different social media platforms, languages, and domains is essential for real-world applications. Incorporating contextual understanding and domain knowledge into deep learning architectures can further enhance performance and address challenges such as context-dependent sarcasm detection. Future research directions may involve exploring multimodal approaches, combining text with other modalities such as images or user metadata, to improve sentiment analysis and sarcasm detection in social media data. Continued research efforts are needed to address these challenges and advance the effectiveness of deep learning approaches in sentiment analysis and sarcasm detection task

CHAPTER 3

METHODOLOGY

3.1 DATA COLLECTION:

The "Large Movie Review Dataset integrated with Keras" is a curated collection of movie reviews specifically designed for sentiment analysis tasks, focusing on IMDb movie reviews.

The dataset contains a substantial number of movie reviews sourced from IMDb, one of the most extensive online databases for films and television shows. This large collection ensures that there is an ample amount of data available for training and testing sentiment analysis models. IMDb serves as the primary source of the movie reviews included in the dataset. IMDb is a reputable platform where users can submit reviews and ratings for movies, providing a diverse range of opinions and sentiments expressed by moviegoers. The dataset likely encompasses a diverse range of opinions and sentiments expressed by users on IMDb. This diversity ensures that the sentiment analysis models trained on the dataset can generalize well to various types of reviews, including both positive and negative sentiments.

Each review in the dataset is manually labeled with its sentiment, indicating whether it expresses a positive or negative opinion about the movie. This labeling is crucial for supervised learning, where the sentiment analysis models learn to associate specific textual patterns with positive or negative sentiments. Human annotators read each movie review and determine whether the overall sentiment conveyed in the review is positive or negative. This manual process requires human judgment and understanding of the nuances of language to accurately interpret the sentiment expressed in the text. The sentiment labels assigned to the reviews are typically binary, meaning each review is labeled as either positive or negative. This binary classification simplifies the sentiment analysis task by categorizing reviews into two distinct sentiment categories. The labeled dataset serves as the training data for supervised learning algorithms in sentiment analysis. In supervised learning, the sentiment analysis models learn to associate specific textual patterns or features with the sentiment labels provided in the training data. By training on the labeled dataset, the sentiment analysis models learn to recognize and extract textual features indicative of positive or negative sentiments. These features may include words, phrases, sentiment expressions, or linguistic cues that signal positive or negative opinions in the reviews. The goal of supervised learning is to enable the sentiment analysis models to generalize their learned associations beyond the training data. Once trained on the labeled dataset, the models can effectively analyze new, unseen movie reviews and predict their sentiments based on the learned patterns. Accurate and reliable

labeling of the dataset is crucial for training high-performing sentiment analysis models. Consistent and well-defined labeling standards help ensure that the models learn to make accurate predictions about the sentiments expressed in movie reviews. The dataset includes movie reviews from a wide array of films, spanning different genres, languages, release years, and cultural backgrounds. It encompasses movies from various genres such as action, drama, comedy, romance, thriller, horror, and more. By covering a diverse range of movies, the dataset ensures that sentiment analysis models are exposed to a broad spectrum of content, reflecting the rich diversity of cinematic experiences.

The diverse coverage of movies and genres enables sentiment analysis models to generalize well to different types of movie reviews. Models trained on a comprehensive dataset are better equipped to handle the variability in language, themes, and stylistic elements present across diverse films and genres. Generalization ensures that sentiment analysis models can accurately predict sentiments for a broad range of movie reviews, regardless of their specific characteristics or subject matter.

Exposure to a diverse range of movies and genres helps in creating sentiment analysis models that are robust and adaptable. Models trained on a comprehensive dataset are less likely to be biased or limited in their predictive capabilities by specific genres or movie types. The robustness of the models allows them to accurately analyze sentiments across different films and genres, yielding more reliable predictions and insights.

In real-world scenarios, movie reviews can vary widely in content, tone, and sentiment, reflecting the diverse preferences and opinions of audiences. Sentiment analysis models trained on a comprehensive dataset are better equipped to handle the variability encountered in real-world movie review data, making them more applicable and effective in practical settings. The ability to generalize across diverse movies and genres enhances the models' utility in analyzing sentiment trends, identifying audience preferences, and informing decision-making processes in the film industry.

Balanced datasets contain an equal distribution of instances across different classes or labels, in this case, positive and negative sentiments. In sentiment analysis, a balanced dataset ensures that the model is exposed to an equal number of examples representing each sentiment category during training. By maintaining balance, the model is less likely to be biased towards predicting one sentiment over the other, resulting in more accurate and reliable predictions across all sentiment classes. Balanced datasets help in training models that can generalize well to diverse sentiment distributions encountered in real-world data.

The dataset is divided into training and testing sets. The training set is used to train the sentiment analysis models, while the testing set is used to evaluate their performance on

unseen data. This ensures that the models' performance can be assessed objectively and that they can generalize well to new, unseen movie reviews.

The training set comprises a subset of the dataset used to train the sentiment analysis models. During training, the models learn to recognize patterns and associations between textual features and sentiment labels present in the training data. The training process involves adjusting the model's parameters through iterative optimization techniques, such as gradient descent, to minimize the error in predicting sentiments on the training set. The testing set is a separate subset of the dataset that is kept unseen by the model during training. Once the model is trained, it is evaluated on the testing set to assess its performance on new, unseen data. The testing process involves feeding the testing set into the trained model and measuring its ability to accurately predict sentiments for the reviews in the testing set. The performance metrics, such as accuracy, precision, recall, and F1-score, calculated on the testing set provide insights into the model's generalization and predictive capabilities.

Splitting the dataset into training and testing sets ensures an objective assessment of the model's performance. By evaluating the model on unseen data from the testing set, researchers can obtain an unbiased estimate of its ability to generalize to new, unseen movie reviews. This objective evaluation helps identify any potential issues, such as overfitting or underfitting, and allows for adjustments to the model's architecture or hyperparameters to improve its performance.

The testing set simulates real-world scenarios where the model encounters new, unseen movie reviews that were not part of the training data. A model that performs well on the testing set demonstrates its ability to generalize effectively to new data, including reviews for movies not present in the training set. Generalization is crucial for ensuring the model's utility in practical applications, where it may need to analyze sentiments for a wide range of movie reviews beyond those included in the training data.

In addition to a single train-test split, researchers may employ techniques like k-fold cross-validation to further validate the model's performance. Cross-validation involves splitting the dataset into multiple subsets (folds), training the model on a combination of folds, and evaluating its performance on the remaining fold(s). This iterative process provides a more robust estimate of the model's performance by averaging results across multiple train-test .

3.2. DATA PREPROCESSING:

Before training the sentiment analysis models, the dataset undergoes pre-processing techniques such as tokenization and word embedding and potentially other techniques to prepare the raw text data for sentiment analysis, enhance model performance, and provide valuable insights into audience sentiments towards IMDb movie reviews.

Tokenization involves breaking down the raw text of IMDb movie reviews into smaller, meaningful units called tokens, typically words or subwords. Word embedding is a technique used to represent words as dense vectors in a continuous vector space. Each word in the vocabulary is mapped to a high-dimensional vector, where similar words are represented by vectors that are close together in the vector space. These preprocessing steps enable the sentiment analysis models to understand the semantic meaning and contextual relationships between words in the reviews. By representing words as dense vectors, word embedding helps address the sparsity and high-dimensionality issues associated with traditional one-hot encoding techniques. Importance of Preprocessing Techniques: Preprocessing techniques such as tokenization and word embedding are crucial for preparing the raw text data for sentiment analysis. Tokenization standardizes the input data by breaking it down into smaller units, facilitating further analysis and feature extraction. Word embedding captures the semantic meaning and contextual relationships between words, enabling the models to understand the underlying meaning of the text and make more informed predictions about sentiment.

Word embedding allows sentiment analysis models to capture the semantic meaning and contextual relationships between words in the reviews. By representing words as dense vectors, word embedding helps the models understand the nuanced meanings of words within the context of the reviews and make more accurate predictions about sentiment.

Preprocessing techniques such as tokenization and word embedding help improve the performance of sentiment analysis models by providing them with richer, more meaningful representations of the input data. Models trained on tokenized and embedded text data are better equipped to capture the subtle nuances of language and make accurate predictions about sentiment.

It also mentions the importance of other preprocessing techniques, such as attention mechanisms and ensemble learning methods, which could further enhance the robustness and generalization capabilities of the sentiment analysis system. These additional preprocessing techniques focus on identifying the most relevant parts of the reviews and incorporating multiple models to improve overall performance.

3.3 PROJECT PHASE :

Data Collection and Preprocessing, the focus is on gathering a sizable dataset of text data with associated sentiment labels and preparing this data for use in training sentiment analysis models. The first step is to collect a substantial amount of text data containing IMDb movie reviews. This dataset serves as the foundation for training and evaluating sentiment analysis models. The dataset should cover a diverse range of movies and genres to ensure that the sentiment analysis models generalize well to various types of movie reviews. Platforms like IMDb provide valuable sources for obtaining movie reviews, which can be scraped or accessed through APIs to gather the required data.

Along with the text data, sentiment labels are required to indicate whether each movie review expresses a positive, negative, or neutral sentiment. These sentiment labels can be obtained through manual annotation by human annotators or may already be available as part of the dataset.

Data preprocessing involves several steps to clean and prepare the raw text data for analysis: Cleaning: Remove any irrelevant or noisy elements from the text data, such as HTML tags, special characters, punctuation, and stop words. Tokenization: Break down the cleaned text data into smaller units called tokens, typically words or subwords. This step facilitates further analysis and feature extraction. Vectorization (Possibly): Depending on the chosen modeling approach, the text data may need to be converted into numerical vectors for model input. Techniques such as word embedding or one-hot encoding can be used for this purpose. Normalization: Standardize the text data by converting all letters to lowercase to ensure consistency in text representation.

It's crucial to ensure that the dataset has a balanced representation of positive and negative sentiment labels to prevent bias in model training. Techniques such as stratified sampling or data augmentation may be employed to achieve a balanced distribution of sentiment labels within the dataset.

Finally, perform quality checks to ensure the integrity and reliability of the collected dataset. This may involve manual inspection of a subset of the data to verify the correctness of sentiment labels and data cleanliness. The dataset is prepared and ready for use in training sentiment analysis models. The cleaned and preprocessed text data, along with associated sentiment labels, serve as the input for subsequent phases of model development, evaluation, and optimization.

Model Selection and Development, the focus is on experimenting with different neural network architectures, such as CNNs and LSTMs, to determine the most suitable models for sentiment analysis of IMDB movie reviews.

The first step is to explore and experiment with various neural network architectures commonly used for natural language processing (NLP) tasks, particularly for sentiment analysis. This may include architectures such as Convolutional Neural Networks (CNNs), Long Short-Term Memory networks (LSTMs), or their combinations like CNN- LSTMs. Each architecture has its strengths and weaknesses, and experimenting with multiple architectures helps identify the most effective approach for the specific task of sentiment analysis on IMDB movie reviews.

Once the architectures are selected, the next step is to train the chosen models on the preprocessed dataset of IMDB movie reviews. During training, the models learn to extract features from the input text data and make predictions about the sentiment expressed in the reviews. Training involves iterative optimization of model parameters, typically using techniques such as gradient descent and backpropagation to minimize the loss function and improve model performance.

Hyperparameters are parameters that are set before the training process begins and control the learning process of the model. Experimentation with different hyperparameter configurations, such as learning rate, batch size, dropout rate, and number of layers, can significantly impact model performance. Hyperparameter tuning involves systematically exploring different combinations of hyperparameters to find the optimal configuration that maximizes model performance.

Cross-validation techniques, such as k-fold cross-validation, may be employed to assess the robustness of the trained models and ensure their generalization to unseen data. Cross-validation involves splitting the dataset into multiple subsets (folds), training the models on a combination of folds, and evaluating their performance on the remaining fold(s). This process is repeated multiple times to obtain a more reliable estimate of model performance.

After training, evaluate the performance of the trained models using evaluation metrics such as accuracy, precision, recall, F1-score, and area under the receiver operating characteristic curve (AUC- ROC). Compare the performance of different models to identify the most effective architecture for sentiment analysis of IMDB movie reviews.

The model development process may involve iterative refinement, where adjustments are made to the architectures, hyperparameters, or preprocessing techniques based on insights gained from

performance evaluation. Continuously iterate and refine the models to improve their accuracy, robustness, and generalization capabilities. The most suitable neural network architectures for sentiment analysis of IMDB movie reviews are identified, trained, and evaluated. These models serve as the basis for further optimization and integration in subsequent phases of the project.

Hyperparameter Tuning and Optimization, the objective is to fine-tune the selected model's hyperparameters and explore techniques like regularization and dropout to enhance its performance. **Hyperparameter Tuning:** Hyperparameters are parameters that are not learned during the training process but are set before training begins. They control the learning process of the model and can significantly impact its performance. Hyperparameter tuning involves systematically searching for the optimal values of these parameters to improve the model's performance. Techniques such as grid search, random search, or more advanced optimization algorithms like Bayesian optimization can be used to explore the hyperparameter space and find the best configuration.

Regularization techniques are used to prevent overfitting, which occurs when the model learns to memorize the training data instead of generalizing well to unseen data. Regularization helps prevent the model from becoming too complex and ensures that it generalizes well to new, unseen data by reducing the variance in the model's predictions.

Dropout is a regularization technique specific to neural networks that helps prevent overfitting by randomly dropping out (setting to zero) a proportion of neurons during training.

After fine-tuning the hyperparameters and applying regularization techniques, evaluate the performance of the optimized model using appropriate evaluation metrics. Compare the performance of the optimized model to the baseline model and previous iterations to assess the effectiveness of the hyperparameter tuning and regularization techniques. The selected model is fine-tuned and optimized to improve its performance and prevent overfitting. The optimized model serves as the final version for deployment and integration into real-world applications for sentiment analysis of IMDB movie reviews.

The system is configured to accept input movie reviews from users or other sources, which are then passed to the sentiment analysis model for prediction. Input handling mechanisms are implemented to preprocess the input data, tokenize it, and prepare it for input to the model. The sentiment analysis model predicts the sentiment of the input movie reviews as either positive or negative. The predicted sentiments are then communicated back to the user or incorporated

into the target application or platform, such as displaying sentiment analysis scores alongside movie reviews.

Validation and Testing, the objective is to validate the trained model using cross-validation techniques to ensure its robustness and then test the final model on unseen data to assess its real-world performance.

Cross-validation is a technique used to assess the robustness and generalization performance of the trained model. The dataset is divided into multiple subsets or folds, typically k folds, where each fold is used as a validation set while the remaining folds are used for training. The model is trained k times, each time using a different fold as the validation set and the remaining folds for training. Cross-validation helps estimate the variability of the model's performance and provides insights into its stability across different subsets of the data.

After cross-validation, evaluate the performance of the trained model using appropriate evaluation metrics such as accuracy, precision, recall, F1-score, and area under the receiver operating characteristic curve (AUC-ROC). Compare the performance of the model across different folds to assess its robustness and generalization capabilities.

Once the model has been validated using cross-validation, it is tested on unseen data to assess its real-world performance. Unseen data typically consists of a separate test set that was not used during training or validation. The model's performance is evaluated on the test set using the same evaluation metrics used during cross-validation.

Evaluate the performance of the final model on the unseen test data to assess its ability to generalize to new, unseen IMDB movie reviews. Compare the performance of the final model to the validation results obtained during cross-validation to ensure consistency and reliability.

If the performance of the final model on the test data is not satisfactory, further iterations of model refinement and optimization may be necessary. This may involve revisiting earlier phases of the project, such as model selection and hyperparameter tuning, to improve the model's performance. By completing Phase 4, the trained model is validated using cross-validation techniques to ensure its robustness and then tested on unseen data to assess its real-world performance. The evaluation results obtained during this phase provide insights into the model's generalization capabilities and help determine its readiness for deployment in real-world applications for sentiment analysis of IMDB movie reviews.

Integration and Deployment, the trained sentiment analysis model is integrated into an application or platform where sentiment analysis is needed, such as a movie review website or a recommendation system.

The trained sentiment analysis model is integrated into the target application or platform, ensuring seamless communication between the model and other components of the system. Depending on the deployment environment, the model may be integrated as part of a web service, API, or a standalone application.

The integrated system is designed to be scalable and efficient, capable of handling varying inputs and processing large volumes of data efficiently. Techniques such as load balancing, parallel processing, and distributed computing may be employed to ensure optimal performance and scalability of the system. Once deployed, the system is continuously monitored to ensure its performance meets the required standards. Performance metrics such as response time, throughput, and accuracy are monitored, and any deviations from expected behavior are addressed promptly. Regular maintenance and updates may be necessary to address issues, improve performance, and incorporate new features or improvements to the sentiment analysis model.

The trained sentiment analysis model is successfully integrated into an application or platform where sentiment analysis is needed. The deployed system is scalable, efficient, and capable of handling varying inputs, providing users with accurate predictions of movie review sentiments in real-time. Ongoing monitoring, maintenance, and iterative improvements ensure the continued effectiveness and reliability of the sentiment analysis system.

CHAPTER 4

SYSTEM DESIGN

4.1 ARCHITECTURE

4.1.1 Architecture of LSTM




Figure 1: Architecture of L S T M

The input layer in the LSTM models, represents the initial stage where the data is fed into the network for processing. In the context of sentiment analysis of IMDb movie reviews, the input layer represents the movie review text, which needs to be transformed into a format suitable for processing by the neural network.

In sentiment analysis, IMDb movie reviews are often represented as sequences of words or tokens. Each word in the review needs to be converted into a numerical representation that the neural network can understand and process. This conversion is typically done using techniques like word embeddings. Word embeddings are dense, lower-dimensional representations of words that capture semantic relationships between words based on their context in the training data.

Word embeddings map each word in the vocabulary to a high-dimensional vector in a continuous space. Words with similar meanings or contexts are mapped to similar vectors, allowing the neural network to capture semantic relationships between words. These word embeddings are often pre-trained on large text corpora using algorithms like Word2Vec, GloVe, or FastText.

Once the word embeddings are obtained, the input layer represents the movie review text as a sequence of word embeddings. Each word in the review is replaced by its corresponding word embedding vector. The resulting input sequence is a series of word embeddings, where each embedding vector represents a word in the movie review.

The input layer typically expects the input data to be in the form of a matrix, where each row represents a word embedding vector. If the movie review consists of a sequence of N words, and each word is represented by a word embedding of size M , then the input matrix will have dimensions $N \times M$. Once the input sequence is transformed into the appropriate format, it is passed to the next

layer in the neural network architecture, which could be an LSTM layer, a bidirectional LSTM layer, or any other type of layer depending on the model architecture. In summary, the input layer in sentiment analysis models for IMDb movie reviews takes the raw text of the movie review, converts it into word embeddings, and represents it as a sequence of word embeddings, which is then fed into the subsequent layers of the neural network for further processing and analysis.

An LSTM (Long Short-Term Memory) layer is a type of recurrent neural network (RNN) layer commonly used in deep learning for sequential data processing tasks such as natural language processing, time series prediction, and speech recognition.

A forward LSTM layer is a specific configuration of an LSTM layer where the input sequence is processed in the forward direction, meaning that the LSTM units process the input data from the beginning of the sequence to the end.

The input to the forward LSTM layer is a sequence of vectors or tensors. Each element of the sequence corresponds to one time step in the input data. Forward Processing: At each time step, the forward LSTM units process the input data and update their internal states. These states include the cell state (memory) and the hidden state. The hidden state contains information about the current time step, while the cell state retains information over longer sequences.

LSTM units have gate mechanisms, including input gates, forget gates, and output gates. These gates regulate the flow of information through the cell state, allowing the LSTM to selectively retain or forget information from previous time steps.

The output of the forward LSTM layer typically consists of the final hidden state at each time step or the final hidden state of the last time step. This output can be further processed by additional layers in the neural network or used for tasks such as classification, regression, or sequence generation. By processing input sequences in the forward direction, forward LSTM layers can capture temporal dependencies and patterns in the data, making them effective for tasks where the order of elements in the sequence matters.

A backward LSTM layer is similar to a forward LSTM layer but processes the input sequence in the opposite direction. While a forward LSTM layer processes the input sequence from the beginning to the end, a backward LSTM layer processes it from the end to the beginning.

Like a forward LSTM layer, the input to a backward LSTM layer is a sequence of vectors or tensors. Each element of the sequence corresponds to one time step in the input data.

At each time step, the backward LSTM units process the input data in reverse order, starting from the end of the sequence and moving towards the beginning. Hidden States and Cell States: Similar to forward LSTM units, the backward LSTM units maintain hidden states and cell states. These states are updated at each time step based on the input and the previous states.

Backward LSTM units also have gate mechanisms, including input gates, forget gates, and output gates. These gates regulate the flow of information through the cell state, allowing the LSTM to selectively retain or forget information from previous time steps, but in this case, moving backward in time.

The output of the backward LSTM layer typically consists of the final hidden state at each time step or the final hidden state of the first time step (which corresponds to the end of the input sequence).

By processing input sequences in the backward direction, backward LSTM layers can capture different temporal dependencies and patterns in the data compared to forward LSTM layers. In some cases, using both forward and backward LSTM layers together in a bidirectional LSTM can be beneficial for capturing comprehensive information from input sequences.

A concatenation layer, often used in the context of bidirectional recurrent neural networks (RNNs), combines the outputs from the forward and backward LSTM (or any other type of RNN). By incorporating Dropout layers into a neural network architecture, you can effectively reduce overfitting and improve the generalization performance of the model on unseen data. However, it's essential to tune the dropout rate appropriately, as setting it too low may not provide sufficient regularization, while setting it too high may hinder the learning process.

layers at each time step.

The forward LSTM layer processes the input sequence in the forward direction, producing an output at each time step. Each output represents the hidden state of the forward LSTM at that time step.

Similarly, the backward LSTM layer processes the input sequence in the backward direction, generating an output at each time step. Each output corresponds to the hidden state of the backward LSTM at that time step. At each time step, the outputs from the forward and backward LSTM layers are concatenated together. This concatenation operation combines the information captured by both the forward and backward LSTMs into a single representation for that time step. The combined output from the concatenation layer contains information from

both the forward and backward directions of the input sequence. This combined representation can capture a more comprehensive understanding of the input sequence, leveraging information from both past and future contexts.

Typically, the outputs from the forward and backward LSTM layers have the same dimensionality. Therefore, when concatenated, the dimensionality of the combined output is twice that of each individual LSTM output.

The concatenated output can then be passed to subsequent layers in the neural network for further processing, such as classification, regression, or sequence generation. Concatenating the outputs from both directions is a common technique used to improve the performance of RNNs, especially in tasks where capturing bidirectional context is beneficial, such as natural language processing and sequence labeling.

A Dropout layer is a form of regularization used in neural networks to prevent overfitting. Overfitting occurs when a model learns to perform well on the training data but fails to generalize to unseen data, resulting in poor performance on test or validation data.

During training, a Dropout layer randomly sets a fraction of input units to zero with a probability defined

by a dropout rate parameter. This means that the output of the Dropout layer for those units becomes zero. The dropout rate is typically set between 0 and 1, representing the fraction of input units to drop.

Dropout operates by "dropping out," or ignoring, certain units in the neural network with each training iteration. By randomly dropping units, Dropout introduces noise into the network, forcing it to learn more robust and generalizable features rather than relying too heavily on specific units. By randomly dropping units during training, Dropout prevents the network from memorizing the training data or learning complex, spurious patterns that are unique to the training set. Instead, it encourages the network to learn more diverse and representative features that generalize better to unseen data.

During inference (i.e., when making predictions on new data), no units are dropped out. Instead, the outputs of the Dropout layer are scaled by a factor equal to the dropout rate. This scaling ensures that the expected output remains the same as during training, accounting for the dropped units.

Dropout can be applied to various types of layers in a neural network, including fully connected layers, convolutional layers, and recurrent layers. It can also be applied multiple times throughout the network to increase regularization.

A fully connected layer with ReLU activation, often used in neural network architectures, follows the concatenation layer in a bidirectional LSTM setup to further process and extract features

from the combined representations obtained from both the forward and backward LSTM layers. The concatenated output from the concatenation layer contains information from both the forward and backward directions of the input sequence. This combined representation captures a richer understanding of the input data, leveraging both past and future contexts.

The fully connected layer adds a set of learnable weights to the concatenated representation. Each neuron in this layer is connected to every neuron in the concatenated representation. The purpose of this layer is to perform a linear transformation on the input data. Following the linear transformation, a rectified linear unit (ReLU) activation function is applied element-wise to the output of the fully connected layer. ReLU introduces non-linearity to the model by outputting the input directly if it is positive, and zero otherwise. This activation function helps the network learn complex patterns and relationships in the data.

The fully connected layer with ReLU activation serves to further transform and extract features from the concatenated representations. By learning appropriate weights and applying non-linear activation, this layer can capture higher-level features and representations that are useful for the task at hand.

In addition to feature extraction, the fully connected layer may also perform dimensionality reduction, effectively reducing the dimensionality of the input representation to a more manageable size. This can help reduce computational complexity and overfitting, especially in tasks with high-dimensional input data.

The output of the fully connected layer with ReLU activation serves as input to subsequent layers in the neural network, such as additional fully connected layers, softmax layers for classification tasks, or regression layers for regression tasks.

Overall, the fully connected layer with ReLU activation after the concatenation layer plays a crucial role in further processing and extracting meaningful features from the combined representations obtained from bidirectional LSTM layers, ultimately improving the performance of the neural network on various tasks.

Using dropout after a fully connected layer with ReLU activation allows us to regularize the activations of the ReLU units. ReLU activation introduces non-linearity into the network by setting negative values to zero. Dropout applied after ReLU helps in preventing the model from overfitting to the training data, especially when dealing with deep neural networks with multiple layers.

The output layer of a neural network is the final layer that produces the model's predictions or outputs. Its configuration depends on the nature of the task being performed, such as binary classification, multi-class classification, or regression.

For binary classification tasks like sentiment analysis, the output layer typically consists of a single neuron with a sigmoid activation function. The sigmoid activation function squashes the output of the neuron to a range between 0 and 1, representing the probability of the positive class (e.g., positive sentiment) given the input. The output can be interpreted as the likelihood or confidence that the input belongs to the positive class. For instance, if the output is 0.8, it suggests an 80.

For tasks with more than two classes, such as image classification with multiple categories, the output layer typically consists of multiple neurons equal to the number of classes. Each neuron corresponds to a class and produces a score or probability for that class. The softmax activation function is commonly used in the output layer for multi-class classification. It computes the probability distribution over all classes, ensuring that the sum of probabilities for all classes equals one. The output of the softmax layer represents the model's confidence scores for each class. The class with the highest probability is chosen as the predicted class.

In summary, the output layer of a neural network tailors its configuration to the specific task at hand. For binary classification, a single neuron with a sigmoid activation function suffices to output the probability of the positive class. For multi-class classification, multiple neurons with softmax activation provide a probability distribution across all classes, enabling the model to predict the most likely class.

4.1.2 Architecture of CNN




Figure 2: Architecture of CNN

The input layer of a neural network represents the initial stage where the raw input data is fed into the model. Its configuration depends on the nature of the input data and the requirements of the task being performed. Here's an explanation of the input layer, particularly in the context of sequential data processing tasks like natural language processing.

In tasks involving sequential data, such as text data in natural language processing (NLP), the input sequence typically consists of a sequence of tokens (words, characters, etc.). Before feeding the input sequence into the neural network, each token is transformed into a suitable numerical representation. One common approach is to use word embeddings, which map each word to a dense vector in a continuous space. Word embeddings capture semantic relationships between words and enable the model to understand the meaning of words based on their contexts. Other representations, such as character embeddings or positional encodings, can also be used depending on the specific requirements of the task.

The input layer must be configured to accept input data in the form of numerical vectors representing the input sequence. The dimensionality of the input representation depends on factors such as the vocabulary size (for word embeddings) or the size of the character set (for character embeddings). Each token in the input sequence is typically represented by a fixed-length vector, and the entire input sequence is represented as a matrix, where each row corresponds to the embedding vector of a token. Before being fed into the input layer, the input data may undergo preprocessing steps such as tokenization (splitting text into tokens), padding (ensuring all sequences have the same length), and possibly normalization or other forms of data cleaning. Once the input sequence is transformed into numerical representations and preprocessed, it is

passed into the input layer of the neural network. The input layer distributes these numerical representations to the subsequent layers of the network for further processing and learning.

The input layer of a neural network plays a crucial role in representing the raw input data in a suitable numerical format that can be processed by the network. In NLP tasks, this often involves encoding the input sequence as word embeddings or other appropriate representations to capture meaningful semantic information.

Convolutional layers are fundamental building blocks in convolutional neural networks (CNNs), particularly in tasks involving image recognition, but they can also be used in other domains such as natural language processing and signal processing. These layers consist of multiple convolutional filters (also known as kernels) that are applied to the input data.

A convolutional filter is a small matrix of weights that scans across the input data. Each filter is typically small spatially but extends through the full depth of the input data. The size of the filter determines the receptive field, which is the area of the input that the filter "sees" at once. Each filter is initialized with random weights and then learns to detect specific patterns or features in the input data through the training process.

As the filters convolve (slide) across the input data, they perform element-wise multiplication between their weights and the corresponding input values, followed by summation. This process effectively computes a weighted sum of the input values within the receptive field of the filter. By learning appropriate weights during training, each filter becomes specialized in detecting certain patterns or features, such as edges, textures, or higher-level patterns.

Convolutional layers capture local patterns and spatial dependencies in the input data. This means that they are particularly effective at capturing features that are spatially close to each other and preserving the spatial relationships between them. For example, in image data, convolutional layers can detect edges, corners, and textures at different locations in the image, regardless of their absolute positions.

Convolutional layers typically consist of multiple filters, each producing a 2D activation map (also known as a feature map) by convolving with the input data. The depth of the output volume is determined by the number of filters in the layer, and each filter produces a separate channel in the output. By stacking multiple channels, convolutional layers can capture complex and multi-scale features in the input data.

After the convolution operation, a non-linear activation function, such as ReLU (Rectified Linear Unit), is typically applied element-wise to the output feature maps to introduce non-linearity into the network.

The convolutional layers play a crucial role in extracting hierarchical representations of input data, enabling CNNs to learn meaningful features directly from the raw input. Through the use of convolutional filters and feature maps, convolutional layers can effectively capture local

patterns and spatial dependencies, making them highly effective for tasks such as image recognition, object detection, and semantic segmentation.

Pooling layers are an essential component of convolutional neural networks (CNNs), often used to down sample the feature maps produced by convolutional layers. These layers help reduce the spatial dimensions (width and height) of the feature maps while retaining important information. Pooling layers reduce the spatial dimensions of the input feature maps by applying a pooling operation over local regions. The most common pooling operations are max pooling and average pooling. Max pooling takes the maximum value from each local region, while average pooling takes the average value.

Pooling is performed independently on each feature map/channel of the input. A pooling window (typically a small square or rectangular region) slides over the input feature map with a specified stride, moving by a certain number of pixels at each step. At each position of the window, the pooling operation (max or average) is applied to the values within the window to produce a single output value for that region.

In addition to spatial dimension reduction, pooling layers also reduce the depth (number of channels) of the feature maps, albeit indirectly. Since each pooling operation produces a single output value for each region, the output feature map will have fewer channels than the input. Despite reducing the spatial dimensions, pooling layers aim to retain the most important features of the input. Max pooling, in particular, retains the most prominent features within each region, while average pooling preserves more global information about the input. This feature retention helps maintain the network's ability to learn hierarchical representations of the input data. Reducing Computational Complexity: Pooling layers help reduce the computational complexity of the network by decreasing the number of parameters and computations required in subsequent layers.

Dropout regularization is a technique commonly used in neural networks to prevent overfitting, which occurs when a model learns to perform well on the training data but fails to generalize to unseen data. Dropout layers are introduced after fully connected layers to implement this regularization technique.

Dropout layers randomly "drop out" (set to zero) a fraction of input units during training. The dropout rate, typically set between 0 and 1, determines the proportion of units to drop. By randomly dropping units, dropout introduces noise into the network and prevents it from relying too heavily on specific features or units, thereby encouraging more robust and generalized learning. Dropout regularization helps prevent overfitting by preventing the network from memorizing the training data or learning spurious correlations that are specific to the training set. By forcing the network to adapt to the absence of certain units during training, dropout encourages the network to learn more diverse and representative features that generalize

better to unseen data.

Dropout layers are typically added after fully connected layers because fully connected layers tend to have a large number of parameters, making them prone to overfitting, especially in deep neural networks. By applying dropout after fully connected layers, dropout regularization can effectively regularize the connections between the layers, reducing the risk of overfitting. Dropout layers can be added to various parts of the neural network, including convolutional layers, recurrent layers, and fully connected layers, depending on the architecture and requirements of the model. While dropout is commonly used after fully connected layers, it can also be applied before or after other types of layers to achieve regularization.

It's important to note that dropout is only applied during training and is typically turned off during inference or when making predictions on new data. During inference, the dropout layers are deactivated, and the output of the network is scaled to compensate for the dropped units, ensuring consistent behavior between training and inference.

Dropout layers added after fully connected layers introduce dropout regularization to prevent overfitting by randomly dropping input units during training. This regularization technique helps improve the generalization performance of the neural network on unseen data, making it more robust and effective for various machine learning tasks.

Sentiment analysis often involves binary classification (positive or negative sentiment), the output layer usually consists of a single neuron. Each neuron in the output layer represents the likelihood or confidence of the input data belonging to the positive sentiment class. The activation function commonly used in the output layer for binary classification tasks like sentiment analysis is the sigmoid function (or logistic function). The sigmoid function squashes the output value to the range $[0, 1]$, representing the probability of the input being positive sentiment. A threshold (typically 0.5) is then applied to interpret the prediction: values above the threshold indicate positive sentiment, while values below indicate negative sentiment.

During inference, after the input data has been processed through the CNN, the output of the output layer neuron(s) is obtained. If the output value is greater than the threshold (0.5), the input data is classified as having positive sentiment. Conversely, if the output value is less than the threshold, the input data is classified as having negative sentiment.

During training, the CNN learns to make accurate predictions by adjusting its weights and biases based on the provided training data and associated labels (positive or negative sentiment). The model's parameters are optimized to minimize a suitable loss function (e.g., binary cross-entropy) that measures the difference between the predicted and actual labels.

The output layer of a CNN used for sentiment analysis contains one neuron that produces predictions indicating whether the input data is classified as positive or negative sentiment.

4.2 DESIGN :

4.2.1 Data Collection and Preprocessing:




Figure 3: Data Collection and Preprocessing

Gathering a sizable dataset of text data with associated sentiment labels involves collecting a large number of movie reviews from IMDb. These reviews should cover a diverse range of movies and genres to ensure the dataset's representativeness and generalizability. Once collected, each review needs to be labeled with its corresponding sentiment, indicating whether it expresses positive, negative. This labeling process can be done manually by human annotators or by utilizing existing labeled datasets.

After collecting the data and labeling it with sentiment labels, the next step is data preprocessing. This phase involves several steps to clean and prepare the raw text data for analysis. The preprocessing steps typically include:

Cleaning: Removing any irrelevant or noisy elements from the text data, such as HTML tags, special characters, punctuation, and stop words. This step helps in reducing noise and improving the quality of the dataset.

Tokenization: Breaking down the cleaned text data into smaller units called tokens, typically words or sub words. Tokenization facilitates further analysis and feature extraction by segmenting the text into meaningful units.

Stemming: Stemming involves reducing words to their root or stem form by removing suffixes. For example, "running" and "runs" would both be reduced to "run." This process helps in consolidating words with similar meanings, thus reducing the vocabulary size and improving model performance.

Lemmatization: Lemmatization, on the other hand, involves reducing words to their base or dictionary form, known as lemma. Unlike stemming, lemmatization considers the context of the word and ensures that the resulting lemma is a valid word. For example, "running" and "ran" would both be lemmatized to "run". Lemmatization tends to be more precise than stemming but can be computationally more expensive.

Overall, data preprocessing plays a crucial role in preparing the text data for sentiment analysis models. By cleaning, tokenizing, and possibly vectorizing the data, we create a structured and standardized dataset that can be effectively utilized for training and evaluating sentiment analysis models, such as L S T M (Long Short-Term Memory) or CNN (Convolutional Neural Network). These models can then learn to accurately predict the sentiment of movie reviews based on the preprocessed text data and associated sentiment labels.

4.2.2 Model Training And Evaluation

After the data collection and preprocessing phases, we have a cleaned and labeled dataset containing IMDb movie reviews with associated sentiment labels. This dataset is ready for use in training sentiment analysis models, as it has undergone cleaning, tokenization, and possibly vectorization, ensuring that the text data is standardized and prepared for model input. **Split Dataset into Training and Testing Sets:** The dataset is divided into two subsets: training and testing. The training set is used to train the sentiment analysis models. The models learn from the patterns and relationships present in this data.

The testing set is used to evaluate the final model's performance after training is complete. It provides an unbiased assessment of the model's ability to predict sentiment on unseen data.

Training Set: Used to train the model's parameters through optimization algorithms like back-propagation. Contains a large amount of labeled data to facilitate learning of sentiment patterns. Helps the model learn to generalize well to new, unseen data.

Testing Set: Used to evaluate the final model's performance after training and hyperparameter tuning. Provides an unbiased estimate of the model's ability to generalize to new, unseen data. Helps in assessing the model's accuracy, precision, recall, and other performance metrics.

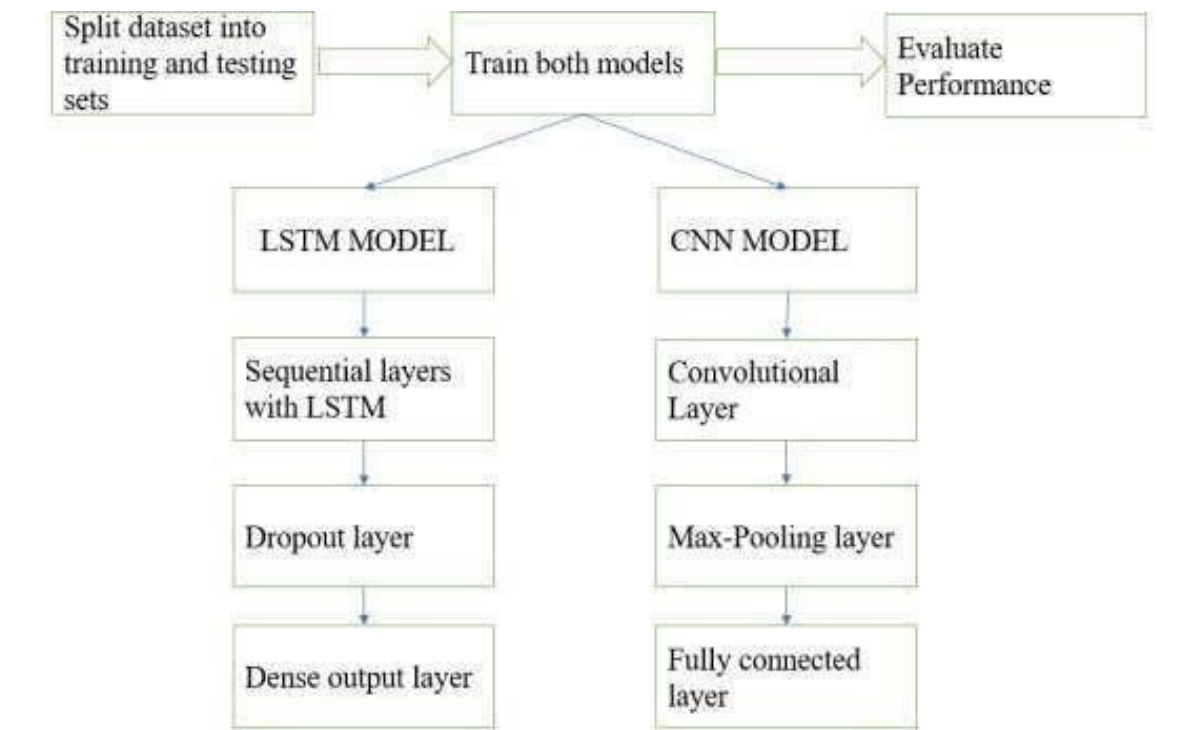


Figure 4: Model Training and Evaluation

[1] Model Building:

LSTM: Long Short-Term Memory networks are well-suited for capturing sequential dependencies in text data. They have memory cells that can store information over time, making them effective for analyzing sequences like sentences or paragraphs. **CNN:** Convolutional Neural Networks excel at capturing local patterns in data. In text analysis, CNNs can learn to recognize important phrases or combinations of words that contribute to sentiment.

Utilize TensorFlow or PyTorch, popular deep learning frameworks, to implement LSTM and CNN architectures for sentiment analysis. These frameworks provide high-level APIs for building neural networks, making it easier to define model architectures and train them efficiently.

Define the number of layers, hidden units, activation functions, and dropout regularization to prevent overfitting. For L STM, configure the number of L S T M layers and the number of units in each layer. Use activation functions like sigmoid or tanh. For CNN, specify the number and size of convolutional filters, pooling layers, and activation functions like ReLU.

[2] Model Training:

Train the L S T M and CNN models using the training dataset. Implement techniques like back-propagation and gradient descent to optimize model parameters and minimize the loss function. Iterate over the training dataset in batches, updating model weights based on the gradient of the loss function with respect to the parameters.

Experiment with different optimization algorithms (e.g., Adam, SGD) and learning rates to find the best combination for training the models. Monitor training progress using metrics like loss and accuracy to ensure that the models are learning effectively. Hyperparameter Tuning:

Use the validation set to tune hyperparameters such as learning rate, batch size, and dropout rate. Evaluate model performance on the validation set after each training epoch to identify the best hyperparameter values.

By systematically building, training, and tuning L S T M and CNN architectures for sentiment analysis using IMDB movie reviews, we can develop robust models capable of accurately predicting sentiment in text data. This iterative process of experimentation and optimization helps in improving model performance and ensuring reliable sentiment analysis results.

[3] Model Evaluation:

Evaluate the trained models using appropriate performance metrics such as accuracy, precision, and recall. These metrics provide insights into the model's ability to classify movie reviews into positive, negative, or neutral sentiments accurately.

Perform cross-validation to assess the model's robustness and generalization across different subsets of the dataset. This technique helps in mitigating the risk of overfitting and provides a more reliable estimate of the model's performance.

Compare the performance of L S T M and CNN models to determine which architecture yields

better results for sentiment analysis of IMDb movie reviews. By rigorously training and evaluating L S T M and CNN models on the IMDb movie review dataset, we can develop robust sentiment analysis models capable of accurately predicting the sentiment expressed in movie reviews across various genres and movies.

4.2.3 Integration and Deployment:

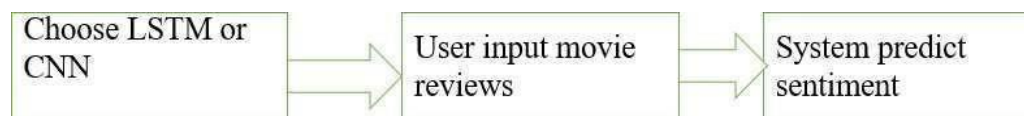


Figure 5: Integration and Deployment

Choose between L S T M and CNN architectures based on performance during model evaluation. Consider factors such as accuracy, computational efficiency, and scalability. L S T M may be preferred if capturing long-range dependencies in text data is crucial, while C N N may excel at capturing local patterns effectively.

Serialize the chosen model (L STM or CNN) along with its trained weights, architecture, and configuration parameters. Deploy the serialized model to a production environment, which could be on cloud platforms like AWS, Google Cloud, or a local server.

Users interact with the system by submitting movie reviews through a user interface or AP I endpoint. The system should handle input validation to ensure that the received text is within acceptable limits and format.

Upon receiving a movie review, the system preprocesses the text data according to the preprocessing steps used during model training. The preprocessed data is fed into the deployed L STM or CNN model for sentiment prediction. The model predicts the sentiment of the movie review as positive, negative, or neutral.

Ensure that the deployed system can handle multiple concurrent user requests efficiently, especially during peak usage times. Monitor system performance and scale resources as needed to maintain responsiveness.

Implement security measures to protect user data and prevent unauthorized access to the system and models. Use encryption for data transmission and storage, and implement authentication and authorization mechanisms.

CHAPTER 5

RESULTS

Collection of IMDb Movie Reviews:

- Testing Data:

TEST DATA

1. Something strange is happening in Loch Ness. The water is crystal clear, sex cold. A giant robotic plastic monster emerges and kills Scotts! What is this movie?! First, I love reading stories about Nessie, sea monsters in general. When I saw this for sale, I thought it was a cheap rip off of jaws. No. It was terrible! The story was pointless, acting was 100% garbage, the only up side was the cool mechanical Nessie they used. It was full of inaccuracy, wrong locations, and had everything. Not worth your while, just leave it on the shelf (or garbage can) you found it on. On second note, This film was shot in California, not Loch Ness, a major dis to Nessie fans.

2. It starts out like a very serious social commentary which quickly makes one think of other Clark movies like Kids, Bully, etc. But then just as quickly, it unravels into a direction-less mess. Who is the main character? Is this a serious film or some Greg Araki-esque over the top goofy film? Is this a skate documentary with moments of dialog inserted? I have no clue. I found myself watching the clock and wonder when this turd was going to end. I kept thinking there would be some big shocker culmination which never came. I cut a good 20 minutes out of the movie by fast forwarding through the pointless skate scenes. Yes, it illustrates the changing landscape between the kids' have no's. I got it way back in the beginning. Kids and Bully was done in such a way that I actually felt like I was observing the realities of that group of friends. Wargup felt very staged, poorly constructed and even worse acting. Teenage Caravan, which Larry didn't write but did direct, was terrible. But at least it felt like it was *supposed* to be a terrible movie that didn't take itself seriously. Wargup Rockers was just plain bad.

3. There's nothing to say except I want my time back that this movie took from me. I'm not racist against Latinos. Hell, I'm half Brazilian. I loved the movie Kids. It doesn't make any sense. These kids just go around and do nothing. They're not even good at skating. The whole time I'm just waiting for something, anything, to happen! but it doesn't. NOTHING happens the whole movie. Did I mention they suck at skating. I might make a movie called Beat up rockers, and the whole premise will be about locking the sh*t out of poorer motion punks like these kids. I'm not even going to get into it, this movie sucks. Please do yourself a favor and burn this movie if you come in contact with it so some other poor soul won't make the same mistake.

4. Ever once in a while I run into a movie that is so embarrassingly bad I wonder why movies exist. This is one of them. This is a terrible attempt to parody The Godfather with annoying cartoon sounds, and bad dialogue. Eddie Deegan is just plain annoying as Tony, an annoying twit who upon his father, Don (William Hickey)'s request, takes over the family business. Tony, as I said, is an annoying little twit. This makes the whole movie a complete mess. The movie is terribly daffy. It's too cartoonish. The main point I'm trying to make is that you can't make a parody of an acclaimed drama like The Godfather with so much cartoonishness. It doesn't work that way. Believe it or not, you have to take a parody of a dramatic movie seriously. If you don't take it seriously, it will feel too much like a parody. The thing about doing a parody is that you can't seem too much like you're doing a parody. You have to make it seem like you're taking the movie at least a little bit seriously. It also feels like they're just mocking Woody Allen, and that's what makes this movie absolutely terrible.

5. Look, don't get me wrong I love independent films but COME ON! I could barely sit through it without wanting to kill myself, the director had absolutely no talent and he turned something that could have been OK to a F***ing Nightmare. I am a punk enthusiast myself and even the music sucked. the acting was *gross*. I am usually a bleeding heart for these low budget films but this one, this one didn't even try. Please don't waste your precious time and money. I am sure to be on hand but comes on! it dragged on and on and on and on. Remember when the kids got into that party with the weird cupcakes and the watermelon martini? how did they just bleed in? It made me frustrated how they could just go anywhere they wanted and get into trouble and have sex and a "meaningful conversation" with whoever they wanted. I know an blathering but my mind is just burning with everything I hated about this film.

6. Great movie. I was laughing all time through. Why? Well, I am from Austria, I can get along with the German (Bavarian) kind of humor. So I guess this movie makes only sense watching when you are German native speaker. Stefan and Erkan both are talking in a new kind of turkish-german accent, which became really popular in our Countries (GER & AUT). But of course they are very stupid. As in every comedy your personal humor will decide, whether thumb up or down.

7. These two main characters Erkan and Stefan are a match comedy act. I was wondering if this is one of these typical slapstick movies where the story is either not important or simply not existing. But when I saw this movie I was very happy that there is a cool story and the main characters really fit in. All in all very amusing and not a common german movie.

8. I enjoyed Erkan & Stefan a cool and fast story which didn't get bogged down in detail. These two guys are great to see and are able to come up with new ideas all the time. The high quality of picture and cut support the movie a lot. Erkan & Stefan show that the German film industry is capable of transferring successful Hollywood concepts with local 'saire' into our cinemas.

9. I was looking in the TV Guide for movies that come from Germany and I found one called The Bunnygrunds, so I watched it and I laughed myself silly! I wanted the DVD but its not available here (I could order it from Germany but it doesn't have subtitles) It was played again so I taped it and watch it from time to time. Anyway, I looked for info on it and found out its real name is Erkan & Stefan, but I know it by its Australian title. The Bunnygrunds. Some people who I know from Germany do not like Erkan & Stefan because of their accents, but not being German myself, I didn't notice anything. The jokes are good, but some German might find their accents off-putting. I think this movie is funny and if the DVD had English subtitles on all the extras (having a 2 disk edition with only the feature having subtitles would be bad) I would buy it up in a snap! I recommend it to anyone looking for a laugh and a pretty good story.

10. quite good, don't expect anything high culture. the acting is bad, the storyline fails, but it is still a fairly nice movie to watch. why? because it's dark, a little bit stupid, like unpredictable and just entertaining and fun to watch. like i said, just see it for yourself and you know what i mean. it is a movie, without a plot or memorable acting, but there are enough scenes that will make you laugh, cry or at least make you feel compelled to watch it to the end. this is all i wanted to say. 7/10

Figure 6: Testing Dataset

• Training Data:

TRAIN DATASET

1. This is an excellent start to the film career of Mickey Rooney. His talents here shows that a long career is ahead for him. The car and truck chase is exciting for the 1937 era. This start of the Andy Hardy series is an American treasure in my book. Spring Byington performance is excellent as usual. Please Mr. Rooney or owners of the film rights, take a chance and get this produced on DVD. I think it would be a winner.

2. Tintin and I recently aired as an episode of PBS's P.O.V. series. It's based on a taped interview of Georges Remi a.k.a. Hergé, Tintin's creator, from 1971 in which he discusses his various experiences publishing his popular character, first in a Catholic newspaper, then in his own series of comic books. Awe-inspiring sweeping views of various comic pages and surreal images of Hergé's dreams. I first encountered Tintin in the pages of Children's Digest at my local elementary school library reading The Secrets of the Unicorn. My mom later got a subscription to CD and I read the entire Rod Kackham's Treasure every month in 1978. I remember seeing some Tintin comic books in a local book store after that but for some reason I didn't get any probably because I was 12 and I thought I was outgrowing them. I do have Breaking Free, a book written and drawn by J. Daniels, published in 1989, six years after Hergé's death. Haven't read it yet. This film also covers the artist's personal life as when he left his first wife after his affair with a colorist in his employ (whom he later married). Her name is Fanny and she is interviewed here. If you love Tintin and his creator, this film is definitely worth a look. Update: 9/4/07-I've now read Breaking Free. Tintin and The Captain are the only regular characters that appear here and they are tailored to the anti-capitalist views of Mr. Daniels with Tintin portrayed as a rabble rouser with a chip on his shoulder who nevertheless cares for The Captain who he's staying with. The Captain here is just trying to make ends meet with a wife and daughter that he loves dearly. They and other construction workers vow to strike after a fellow employee dies from a faulty equipment accident. The whole thing takes place in England with working-class cockney accents intact. Not the kind of thing Hergé would approve of but an interesting read nonetheless. Oh, yes, dog Snowy only appears in the top left corner of the cover (which has Tintin running over the police!) and the dedication page.

3. *Possible Spoilers Ahead*
 > Whenever fans of bad movies congregate for more than a few minutes, a name that invariably comes up is that of Larry Buchanan. This amazing director has given us remakes of other turkeys (ZONTAR THE THING FROM VENUS), cheap-jack crime dramas like A BULLET FOR PRETTY BOY, and tawdry conspiracy flicks like DOWN ON US AND GOODBYE NORMA JEAN. THE LOCH NESS HORROR is a humdinger to say the least. Overlooking the fact that Loch Ness is extremely long and narrow, Larry filmed this boulder on a wide and round California lake. Early on, the film boasts some dazzling (for the budget) underwater photography and creates some atmosphere in spite of itself. Then it degenerates into windy dialogue uttered by no-name actors with lapping Scottish accents, not to mention a soundtrack that will do nothing for the much-maligned bagpipe. At one point, campers sing "You Take The High Road, I'll Take The Low Road," just to throw in one more Scottish cliché. If Scottish people ever decide to jump on the Political Correctness bandwagon they'll sue Larry Buchanan over this film, his surname notwithstanding. The monster looks like a giant papier-mâché puppet and it makes the dragon in Beanie & Cecil look terrifying by comparison. In one unforgettable scene Nessie takes to land and, to evade some patrolling soldiers, the fifty-foot-long critter tries to hide behind a tree and the soldiers don't see it! THE LOCH NESS HORROR is a true mind-boggler that must be seen several times—to be believed.

4. Something strange is happening in Loch Ness. The water is crystal clear, not cold. A giant robotic plastic monster emerges and kills Scots! What is this movie?! First, I love reading stories about Nessie, sea monsters in general. When I saw this for sale, I thought it was a cheap rip off of jaws. No. It was terrible! The story was pointless, acting was 100% garbage, the only up-side was the cool mechanical Nessie they used. It was full of inaccuracy, wrong locations, and bad everything. Not worth your while, just leave it on the shelf (or garbage can) you found it on. On second note, This film was shot in California, not Loch Ness, a major dis to Nessie fans.

5. You get 5 writers together, have each write a different story with a different genre, and then you try to make one movie out of it. If it's action, it's adventure, it's sci-fi, it's western, it's a mess. Sorry, but this movie absolutely stinks. 4.5 is giving it an awfully high rating. That said, it's movies like this that make me think I could write movies, and I can barely write.

6. How wonderful. Yet another movie about America by someone who has visited here probably a half dozen times, a day a piece, and believes himself to be an "expert" on the country. Sheesh. I should take a trip to Germany for a week and then come back and make a movie about Germany as the "land of Nazis" or some such. Wim Il. boy, you should get together with Lars von Trier and make the ULTIMATE movie about the Americans. Of course we all know it takes a pretentious left-leaning "we are the world" European to make a "real" movie about America.
 > Yeah, right. For a continent that started not one but TWO world wars, Europe sure has a lot of opinions about America's wrong "foreign policy".
 > P.S. Don't worry, Wim Il. boy, there's plenty of UC Berkeley Americans that'll just love your movie. Of course, these are the same people who think George W. Bush is worse than Hitler, and that a painting of a can of soup is "sheer genius"!!

7. What happens when you give a free man just enough money to trap him into the rat race and watch him squirm? Homeless people answer to no one. They have no mortgages, rent payments or idiot bosses. Homeless people don't have to worry about the IRS or performance reviews or credit card payments. But, give them just enough money to rent an apartment and buy a car and, suddenly, they have to worry about entering the rat race, buying gas for transportation, paying insurance on their car, and working for someone else. They get a chance to be a "productive citizen." This film was about as exploitive as a film can be. It's a way for the rich and middle-class sheeple to say "see what happens when you try to help the poor?" and it vindicates capitalistic arrogance.
 >
 > Why not a film that asks, "What happens when you take away everything a rich man has?"

8. I've seen this movie and I must say I'm very impressed. There are not much movies I like, but I do like this one. You should see this movie by yourself and comment it because this is one of my most favorite movies. I fancy to see this again. Action fused with a fantastic story. Very impressing. I like Modesty's character. Actually she's very mystic and mysterious (I DO like that!). The bad boy is pretty too. Well, actually this whole movie is rare in 'movieview'. I considered about the vote of this movie, I thought this is should be a very popular movie. I guess wrong. It was ME who was very impressed about this movie, and I hope I'm not the only one who takes only the cost to watch this one. See and vote.

9. The good thing about this film is that it stands alone - you don't have to have seen the original. Unfortunately this is also it's biggest drawback. It would have been nice to have included a few of the original characters in the new story and seen how their lives had developed. Sinclair as in the original is excellent and provides the films best comic moments as he attempts to deal with

Figure 7: Training Dataset

Result after Preprocessing:

- CNN Model:

```
Shape of train_data = (25000, 2494)
Shape of test_data = (25000, 2494)

Sample of preprocessed train_data:
[ 0  0  0 ... 19 178 32]
[ 0  0  0 ... 16 145 95]
[ 0  0  0 ... 7 129 113]
[ 0  0  0 ... 21 64 2574]
[ 0  0  0 ... 7 61 113]
[ 0  0  0 ... 15344 10 10]
[ 0  0  0 ... 55 52 1901]
[ 0  0  0 ... 72 33 32]
[ 0  0  0 ... 28 126 110]
[ 0  0  0 ... 7 43 50]
[ 0  0  0 ... 545 1178 829]
[ 0  0  0 ... 46 7 457]
[ 0  0  0 ... 851 14 2286]
[ 0  0  0 ... 1195 1195 1195]
[ 0  0  0 ... 2 543 1609]
[ 0  0  0 ... 4 2092 151]
[ 0  0  0 ... 358 10397 6]
[ 0  0  0 ... 11 14 2687]
[ 0  0  0 ... 31 33 15]
[ 0  0  0 ... 5 2 25]
```

Figure 8: Training Data of CNN Model

```
406
407 Sample of preprocessed test_data:
408 [ 0  0  0 ... 14 6 717]
409 [ 0  0  0 ... 125 4 3077]
410 [ 0  0  0 ... 9 57 975]
411 [ 0  0  0 ... 28 314 1772]
412 [ 0  0  0 ... 158 856 158]
413 [ 0  0  0 ... 30 23 288]
414 [ 0  0  0 ... 450 5122 596]
415 [ 0  0  0 ... 128 74 14]
416 [ 0  0  0 ... 25 6 3350]
```

Figure 9: Testing Data of CNN Mod

LSTM Model

```
1  Padding the reviews...
2  Shape of train_data = (25000, 2494)
3  Shape of test_data  = (25000, 2494)
4
5  Sample of preprocessed train_data:
6  [ 0 0 0 ... 19 178 32]
7  [ 0 0 0 ... 16 145 95]
8  [ 0 0 0 ... 7 129 113]
9  [ 0 0 0 ... 21 64 2574]
10 [ 0 0 0 ... 7 61 113]
11 [ 0 0 0 ... 15344 10 10]
12 [ 0 0 0 ... 55 52 1901]
13 [ 0 0 0 ... 72 33 32]
14 [ 0 0 0 ... 28 126 110]
15 [ 0 0 0 ... 7 43 50]
16 [ 0 0 0 ... 545 1178 829]
17 [ 0 0 0 ... 46 7 457]
18 [ 0 0 0 ... 851 14 2286]
19 [ 0 0 0 ... 1195 1195 1195]
20 [ 0 0 0 ... 2 543 1609]
21 [ 0 0 0 ... 4 2092 151]
22 [ 0 0 0 ... 358 10397 6]
23 [ 0 0 0 ... 11 14 2687]
```

Figure 10: Training Data of LSTM Model

```
206
207 Sample of preprocessed test_data:
208 [ 0 0 0 ... 14 6 717]
209 [ 0 0 0 ... 125 4 3077]
210 [ 0 0 0 ... 9 57 975]
211 [ 0 0 0 ... 28 314 1772]
212 [ 0 0 0 ... 158 856 158]
213 [ 0 0 0 ... 30 23 288]
214 [ 0 0 0 ... 450 5122 596]
215 [ 0 0 0 ... 128 74 14]
216 [ 0 0 0 ... 45 6 3358]
217 [ 0 0 0 ... 3555 5 2065]
```

Figure 11: Testing Data of LSTM Model

Model Building:

```

from keras.models import Sequential
from keras.layers import Embedding, Conv1D, MaxPooling1D, GlobalMaxPooling1D, Dense, Dropout, Flatten, GlobalAveragePooling1D
from keras.regularizers import l2

model = Sequential()
model.add(Embedding(input_dim = VOCAB_SIZE,
                    output_dim = 128,
                    input_length = max_length))
model.add(Conv1D(filters = 128,
                 kernel_size = 3,
                 strides = 1,
                 padding = 'valid',
                 activation = 'relu',
                 ))
model.add(MaxPooling1D(pool_size = 7))
model.add(GlobalMaxPooling1D())
model.add(Dense(units = 32,
                 activation = 'relu'))
model.add(Dense(units = 1,
                 activation = 'sigmoid'))

model.summary()

```

Figure 12: CNN Model

- **Output of CNN Model:**

```

... Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 2494, 128)	2560000
conv1d (Conv1D)	(None, 2492, 128)	49280
max_pooling1d (MaxPooling1D)	(None, 356, 128)	0
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense (Dense)	(None, 32)	4128
dense_1 (Dense)	(None, 1)	33

```

=====
Total params: 2,613,441
Trainable params: 2,613,441
Non-trainable params: 0

```

Figure 13: CNN Model After Building

```

from keras.models import Sequential
from keras.layers import LSTM, Embedding, Dense, Dropout
from keras.regularizers import l2

model = Sequential()
model.add(Embedding(input_dim=VOCAB_SIZE,
                    output_dim=32,
                    input_length=max_length))
model.add(LSTM(units=128,
               activation = 'tanh',
               recurrent_activation = 'sigmoid',
               dropout = 0.0,
               recurrent_dropout = 0.0,
               return_sequences = False,
               # kernel_regularizer = l2(0.01),
               # recurrent_regularizer = l2(0.01),
               # bias_regularizer = l2(0.01)
               ))
model.add(Dense(units = 64,
                 activation = 'relu'))
model.add(Dropout(0.3))
model.add(Dense(units = 32,
                 activation = 'relu'))
model.add(Dense(units = 1,
                 activation = 'sigmoid'))

model.summary()

```

Figure 14: L S T M Model

- **Output of LSTM Model:**

```

Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 2494, 32)	640000
lstm_2 (LSTM)	(None, 128)	82432
dense_4 (Dense)	(None, 64)	8256
dropout_2 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 32)	2080
dense_6 (Dense)	(None, 1)	33

```

Total params: 732,801
Trainable params: 732,801
Non-trainable params: 0

```

Figure 15: L S T M Model After Building

Model Training:

- CNN Model:

```
history = model.fit(train_data,
                    train_labels,
                    epochs = 100,
                    batch_size = 64,
                    verbose = 1,
                    callbacks = callback_list,
                    validation_split = 0.00512,
                    shuffle = True)
```

Figure 16: CNN Model

- Output of CNN Model

```
389/389 [-----] - 12s 11ms/step - loss: 1.4236e-04 - acc: 1.0000 - val_loss: 0.3574 - val_acc: 0.9062 - lr: 8.0000e-05
Epoch 18/100
389/389 [----->.] - ETA: 0s - loss: 1.3653e-04 - acc: 1.0000
Epoch 18: val_acc did not improve from 0.90625
389/389 [-----] - 12s 11ms/step - loss: 1.3654e-04 - acc: 1.0000 - val_loss: 0.3580 - val_acc: 0.9062 - lr: 8.0000e-05
Epoch 19/100
389/389 [-----] - ETA: 0s - loss: 1.3476e-04 - acc: 1.0000
Epoch 19: val_acc did not improve from 0.90625
389/389 [-----] - 13s 12ms/step - loss: 1.3476e-04 - acc: 1.0000 - val_loss: 0.3585 - val_acc: 0.9062 - lr: 8.0000e-05
Epoch 20/100
389/389 [-----] - ETA: 0s - loss: 1.3256e-04 - acc: 1.0000
Epoch 20: val_acc did not improve from 0.90625
389/389 [-----] - 12s 12ms/step - loss: 1.3266e-04 - acc: 1.0000 - val_loss: 0.3590 - val_acc: 0.9062 - lr: 8.0000e-05
Epoch 21/100
389/389 [-----] - ETA: 0s - loss: 1.3020e-04 - acc: 1.0000
Epoch 21: val_acc did not improve from 0.90625
389/389 [-----] - 12s 10ms/step - loss: 1.3020e-04 - acc: 1.0000 - val_loss: 0.3596 - val_acc: 0.9062 - lr: 8.0000e-05
Epoch 22/100
389/389 [-----] - ETA: 0s - loss: 1.2735e-04 - acc: 1.0000
Epoch 22: val_acc did not improve from 0.90625
Restoring model weights from the end of the best epoch: 2.
Epoch 22: ReduceLROnPlateau reducing learning rate to 1.0000001770593217e-06.
389/389 [-----] - 12s 11ms/step - loss: 1.2735e-04 - acc: 1.0000 - val_loss: 0.3604 - val_acc: 0.9062 - lr: 8.0000e-06
Epoch 22: early stopping
```

Figure 17: CNN Model After Training

- **LSTM Model:**

```
history = model.fit(train_data,
                    train_labels,
                    batch_size = 64,
                    epochs = 100,
                    verbose = 1,
                    callbacks = callback_list,
                    validation_split = 0.00512,
                    shuffle = True)
```

Figure 18: L S T M Model

- **Output of LSTM Model**

[illegible]

Figure 19: L S T M Model After Training

Integration and Deployment:

- **Web Interface:**

IMDb Movie Review Sentiment Analysis

Type a movie review to predict its sentiment.

Enter the movie review:

Predict

Figure 20: Web Interface

- **Output :**

IMDb Movie Review Sentiment Analysis

Type a movie review to predict its sentiment.

Enter the movie review:

Historically, this movie or film/series, is well-regarded by some moviegoers as one of the best that could be regarded for its genius, and is the most famous part of the most famous Western trilogy. Building on a rich Italian storytelling and film making tradition, with unforgettable score, this movie has justly taken the place it now holds. If you have not seen it yet, do it. Equally subversive and authentic, this is just one master piece nobody should miss.

Predict

Prediction (CNN model): Positive

Prediction (LSTM model): Positive

Figure 21: Prediction of the Movie Review

CHAPTER 6

CONCLUSIONS

6.1 CONCLUSIONS

The proposed system aims to improve sentiment analysis of IMDb movie reviews by utilizing advanced neural network architectures within a Python framework. Sentiment analysis involves determining the sentiment or opinion expressed in a piece of text, in this case, movie reviews on IMDb. This task is crucial for various applications like understanding customer feedback, market analysis, and recommendation systems.

Traditional machine learning algorithms may face limitations in handling large datasets efficiently, which is common in platforms like IMDb with a vast number of reviews. To overcome this, the system employs Long Short-Term Memory (LSTM) networks, known for their ability to capture long-range dependencies in sequential data. This makes them suitable for sentiment analysis tasks where the context of the text is crucial.

Additionally, Convolutional Neural Networks (CNNs), typically used in image processing, have shown effectiveness in sequential data analysis, including sentiment analysis. By incorporating both LSTM and CNN architectures separately, the system aims to leverage the strengths of each model. This not only enhances accuracy but also addresses scalability concerns, allowing for efficient processing of large datasets.

Implementing LSTM and CNN architectures separately enables a better understanding of their individual performance characteristics and facilitates fine-tuning of each component. This segregated approach provides a promising strategy for IMDb sentiment analysis, offering improved performance and scalability compared to traditional methods.

6.2 FUTURE WORK:

As for future endeavors, the system suggests exploring cross-domain sentiment analysis models. This entails developing algorithms capable of analyzing sentiment across diverse domains beyond IMDb movie reviews, thereby broadening the applicability and effectiveness of the sentiment analysis framework.

Exploring cross-domain sentiment analysis models offers several benefits: Enhanced versatility: Models trained on diverse datasets can adapt to different types of text data, making them more

versatile in real-world applications. Improved generalization: Models that can generalize sentiment analysis across domains are more robust and reliable, as they are not limited to specific types of text data. Increased insights: Analyzing sentiment across various domains provides a more comprehensive understanding of consumer opinions and attitudes, leading to better-informed decision-making.

REFERENCES

- [1] Vishu Tyagi, Ashwini Kumar, Sanjoy Das, "Sentiment Analysis on Twitter Data Using Deep Learning approach", IEEE'
- [2] Karuna Arava, Rudraraju Sri Krishna Chaitanya, Shaik Sikindar, S Phani Praveen, Swapna D," Sentiment Analysis using deep learning for use in recommendation systems of various public media applications", IEEE.
- [3] Zhang Peiliang, Mijit Ablimit, Hankiz Yilahun, Askar Hamdulla," Dynamic Word Vector Based Sentiment Analysis for Microblog Short Text ", IEEE.
- [4] Nikita Taneja, Hardeo K Thakur," Sentiment-Based Deep Auto Encoder Recommender System (DAERS)", IEEE.
- [5] Aaryan Singh, Harsh Srivastava, Mohd. Aman, Gaurav Dubey, "Sentiment Analysis on User Feedback of a Social Media Platform",IEEE.
- [6] Yi k Yang Tan, Chee-Onn Chow, Jeevan Kanesan, Joon Huang Chuah, YongLiang Lim, "Sentiment Analysis in Social Media: Handling Noisy Data and Detecting Sarcasm Using a Deep Learning Approach", IEEE.
- [7] Jie Yang, Jinxin Yi, Jinming Hu, and Jun Zhang, "A Novel Hybrid Recommendation System Based on Collaborative Filtering and Sentiment Analysis", IEEE.
- [8] Hua Wang, Jianwu Dang, Xiaoming Jin, and Tao Li, "Sentiment-Aware Personalized Recommendation System Using Deep Learning", IEEE.
- [9] Rishabh Gupta, Neha Sharma, Priya Singh, "Enhanced Sentiment Analysis in E-commerce Reviews Using Deep Learning Techniques", IEEE.